

# Kubernetes Workshop

Day 2

# Resource Management

# Requests and Limits

- Requests: Guaranteed resources for scheduling
- Scheduler uses requests to place pods
- Pod gets at least this much
- Limits: Maximum resources allowed
- CPU limit exceeded → Throttled
- Memory limit exceeded → OOMKilled

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  replicas: 3
  template:
    spec:
      containers:
        - name: app
          image: my-app:1.0
          resources:
            requests:
              cpu: "250m"
              memory: "256Mi"
            limits:
              cpu: "1"
              memory: "512Mi"
```

# LimitRange

- Default and max/min constraints per namespace
- Applied to Pods/Containers that don't specify resources

```
apiVersion: v1
kind: LimitRange
metadata:
  name: resource-limits
  namespace: dev
spec:
  limits:
    - type: Container
      default:
        cpu: "500m"
        memory: "256Mi"
      defaultRequest:
        cpu: "100m"
        memory: "128Mi"
      max:
        cpu: "2"
        memory: "2Gi"
      min:
        cpu: "50m"
        memory: "64Mi"
```

# Health Probes

# Health Probes Overview

- Liveness Probe: Is the container alive?
  - Fails → Kubernetes kills and restarts container
  - Use for: detecting deadlocks, hung processes
- Readiness Probe: Is the container ready for traffic?
  - Fails → Removed from Service endpoints (no traffic)
  - Use for: startup warmup, temporary unavailability
- Startup Probe: Has the container started?
  - Disables liveness/readiness until it passes
  - Use for: slow-starting applications

# Probe Types

- httpGet - HTTP request to endpoint
  - Success: HTTP 200-399
  - Best for: Web apps, APIs, microservices
- tcpSocket : TCP connection attempt
  - Success: Connection established
  - Best for: Databases, Redis, message queues
- Exec: Run command in container
  - Success: Exit code 0
  - Best for: Custom health checks

```
apiVersion: v1
kind: Pod
metadata:
  name: probe-demo
spec:
  containers:
  - name: app
    image: my-app:1.0
    ports:
    - containerPort: 8080
  readinessProbe:
    httpGet:
      path: /healthz
      port: 8080
    initialDelaySeconds: 5
    periodSeconds: 10
    failureThreshold: 3
  livenessProbe:
    tcpSocket:
      port: 8080
    initialDelaySeconds: 15
    periodSeconds: 20
    failureThreshold: 3
  startupProbe:
    exec:
      command:
      - cat
      - /tmp/app-initialized
    initialDelaySeconds: 0
    periodSeconds: 5
    failureThreshold: 30
```

# Scheduling and Placement

# NodeSelector

- Simplest way to control pod placement
- Pod only schedules on nodes with ALL matching labels.

```
kubectl label nodes aks-nodepool1-12345 gpu=true disktype:ssd
```

```
apiVersion: v1
kind: Pod
metadata:
  name: gpu-pod
spec:
  nodeSelector:
    gpu: "true"
    disktype: ssd
  containers:
    - name: ml-training
      image: tensorflow/tensorflow:2.20-gpu
```

# Taint

- Blocks all pods by default

```
kubectl taint nodes node1 dedicated=gpu:NoSchedule  
kubectl taint nodes node1 dedicated=gpu:NoSchedule
```

# Tolerations

- Added to pods to allow scheduling on tainted nodes
- Must match key + value + effect
- Toleration does NOT force pod to that node

```
apiVersion: v1
kind: Pod
metadata:
  name: production-app
spec:
  containers:
  - name: app
    image: my-app:1.0

  tolerations:
  - key: "environment"
    operator: "Equal"
    value: "production"
    effect: "NoSchedule"
```

# Node Affinity

- More expressive than nodeSelector

```
apiVersion: v1
kind: Pod
metadata:
  name: required-affinity-pod
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: topology.kubernetes.io/zone
            operator: In
            values:
            - westeurope-1
            - westeurope-2
      containers:
      - name: app
        image: nginx
```

# Pod Affinity and Anti-Affinity

- Pod Affinity: Schedule pods together
  - Use case: Web server + cache on same node (low latency)
- Pod Anti-Affinity: Spread pods apart
  - Use case: Database replicas on different nodes (HA)
- topologyKey defines the domain:
  - kubernetes.io/hostname = same/different node
  - topology.kubernetes.io/zone = same/different AZ

# Pod Disruption Budget

# Pod Disruption Budgets (PDB)

- Protect availability during voluntary disruptions
- Voluntary: node drains, cluster upgrades, autoscaler
- Does not protect against node crashes, pod deletions, app failures
- minAvailable: At least N pods must stay running
- maxUnavailable: At most N pods can be down

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: worker-pdb
spec:
  maxUnavailable: 25%
  selector:
    matchLabels:
      app: worker
```

# Labs

- 2.01: StatefulSets
- 2.02: Requests and Limits
- 2.03: Probes
- 2.04: Scheduling and placement

# Security and RBAC

# Service Accounts


- Identity for pods (not humans)
- Every namespace has a 'default' ServiceAccount
- Pods automatically use 'default' unless specified
- Token mounted at /var/run/secrets/kubernetes.io/serviceaccount/
- Best practice: Create dedicated SA per application
- Don't use 'default' SA for production workloads

# RBAC

- Role-Based Access Control
- Default: Deny everything
- Who? User, Group, or ServiceAccount
- What? Verbs: get, list, watch, create, update, delete
- Where? Resources: pods, deployments, secrets, configmaps
- Role: Permissions in ONE namespace

# RBAC

- Role: Permissions in one namespace
- RoleBinding: Grants Role to subject
- ClusterRole: Permissions cluster-wide
- ClusterRoleBinding: Grants ClusterRole cluster-wide



# Custom Resource Definition

# Custom Resource Definition (CRD)

- Extend Kubernetes API with your own types
- Examples: Database, Certificate, Backup
- Stored and served like built-in resources

# Finalizers

# Finalizers

- Prevents resource deletion until cleanup is done
- Resource stuck in "Terminating" until finalizers removed

# Helm

# Helm

- Package manager for Kubernetes
- Chart: Package of K8s manifests + templates
- Release: Installed instance of a chart
- Repository: Collection of charts
- Values: Configuration for customization
- Templating with Go templates
- Versioning and rollback support

# GitOps

# What is GitOps?


- Git = Single source of truth for infrastructure
- Desired state stored in Git repository
- Automated sync from Git to cluster
- All changes through Git (PRs, reviews, approvals)
- Benefits
  - Audit trail built-in (git history)
  - Rollback
  - Consistency across environments
  - No direct cluster access needed

# ArgoCD

- Declarative, Git-based deployments
- CNCF project
- CLI and web UI
- Monitors Git repos for changes
- Syncs desired state to cluster automatically
- Multi-cluster support



# ArgoCD



# Labs

- 2.05: Service Accounts
- 2.06: Finalizers
- 2.07: Troubleshooting